



# Real-Time Big Data

**Vybornov Artyom**, [avybornov@bigdatateam.org](mailto:avybornov@bigdatateam.org)

Big Data Instructor, <http://bigdatateam.org/>

Head of Big Data Dev Team, Rambler Group

<https://www.linkedin.com/in/artvybor/>



## RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



## Spark Streaming

- ▷ Application example
- ▷ Hints



## Spark Streaming practice



## Kafka

- ▷ Internal
- ▷ CLI



## RT Intro

- ▷ **Batch -> RT**
- ▷ Approaches to RT data processing



## Spark Streaming

- ▷ Application example
- ▷ Hints



## Spark Streaming practice

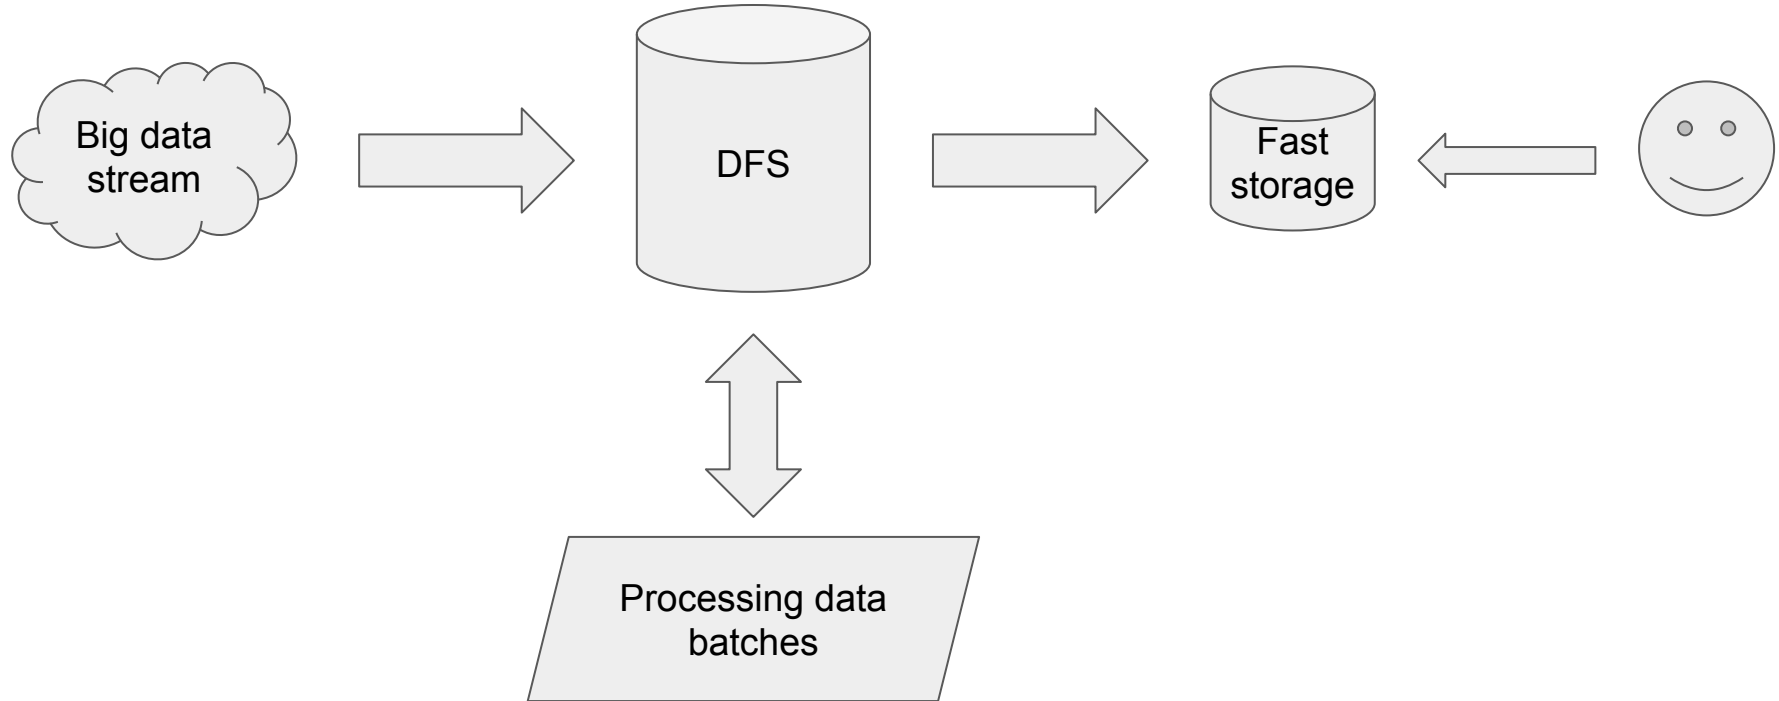


## Kafka

- ▷ Internal
- ▷ CLI



# Batch approach






# The main cons of batch approach

- ▶ In practice a batch is a big time interval like an hour or day
- ▶ The size of batch is a minimal value for a lag
- ▶ Lag (delay) is the time between the event and its recording in the results of work
- ▶ For many purposes the following rule works: less lag => more valuable data



# Real-time big data

- ▶ Real-time big data - is a set of technologies to process big data with minimal lag
- ▶ Without DFS
- ▶ Work with a stream of data instead of a batch





# Real-time big data - lag in minutes

- ▶ Build custom recommendations (Linkedin, Facebook)
  - ▶ Billions of events per day
  - ▶ Millions of events every second
  - ▶ Minutes to build fresh recommendation



# Real-time big data - lag in seconds

## Programmatic advertising (Google, Facebook)

- ▶ Billions of events per day
- ▶ You liked the article and saw the relevant ads already on the next page

## Credit card real-time fraud detection

- ▶ It takes a few seconds to detect a card with suspicious activity and block it





# Real-time big data - lag in ms



## Calls billing (mobile network operator)

- ▶ Hundred millions of users
- ▶ Less than a second to debit funds
- ▶ Zero fault tolerance



## High-frequency trading (HFT)

- ▶ The reaction to the change of quotations of the exchange with a delay of 10-100 milliseconds



## RT Intro

- ▷ Batch -> RT
- ▷ **Approaches to RT data processing**



## Spark Streaming

- ▷ Application example
- ▷ Hints



## Spark Streaming practice

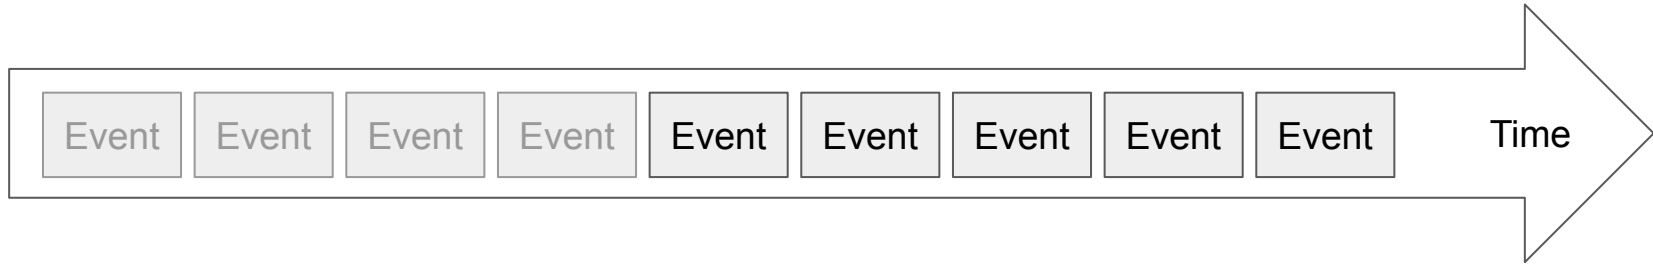


## Kafka

- ▷ Internal
- ▷ CLI

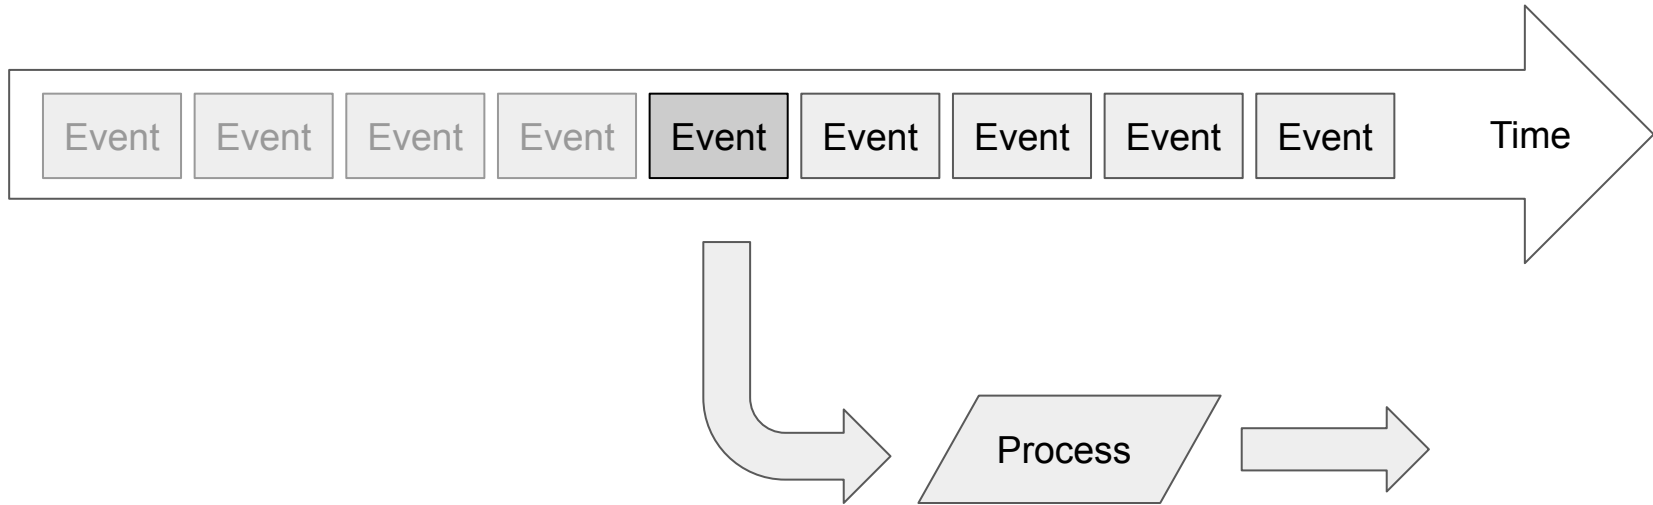


# Event-based approach





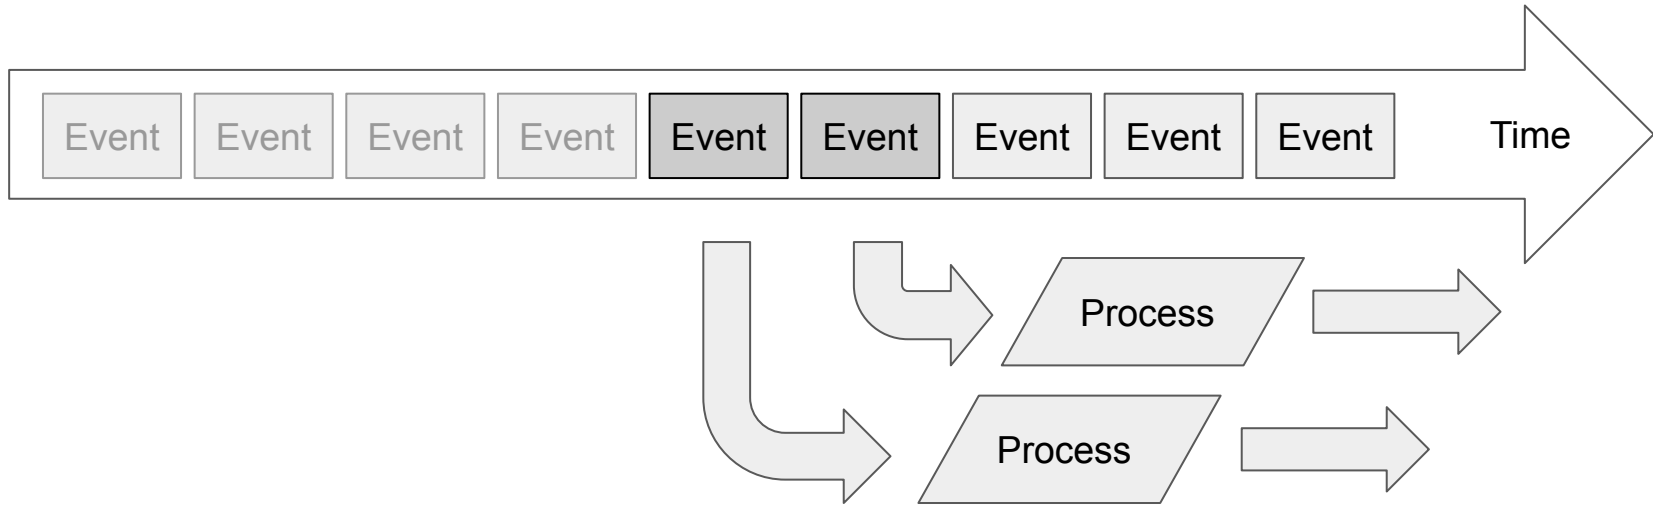
# Event-based approach



Handling one event at a time



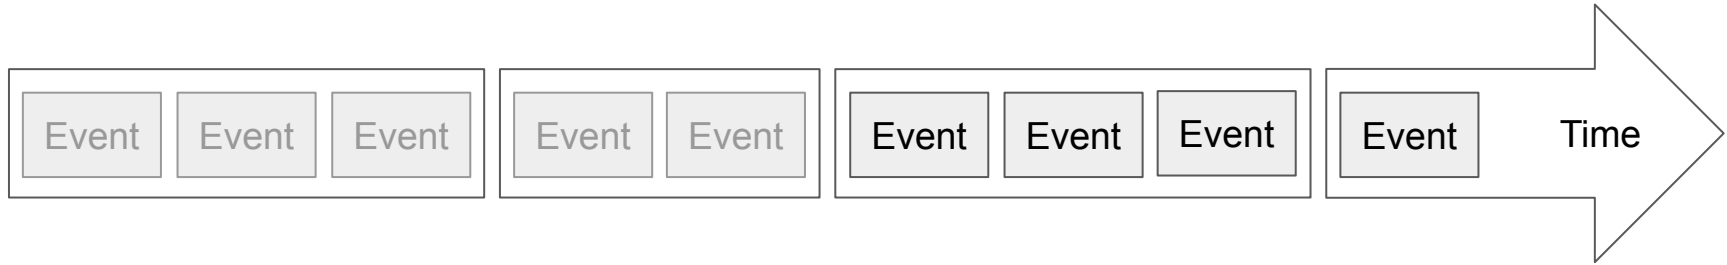
# Event-based approach



- ▶ Handling one event at a time
- ▶ Events are processed in parallel, but completely independently
- ▶ Latency ~10ms



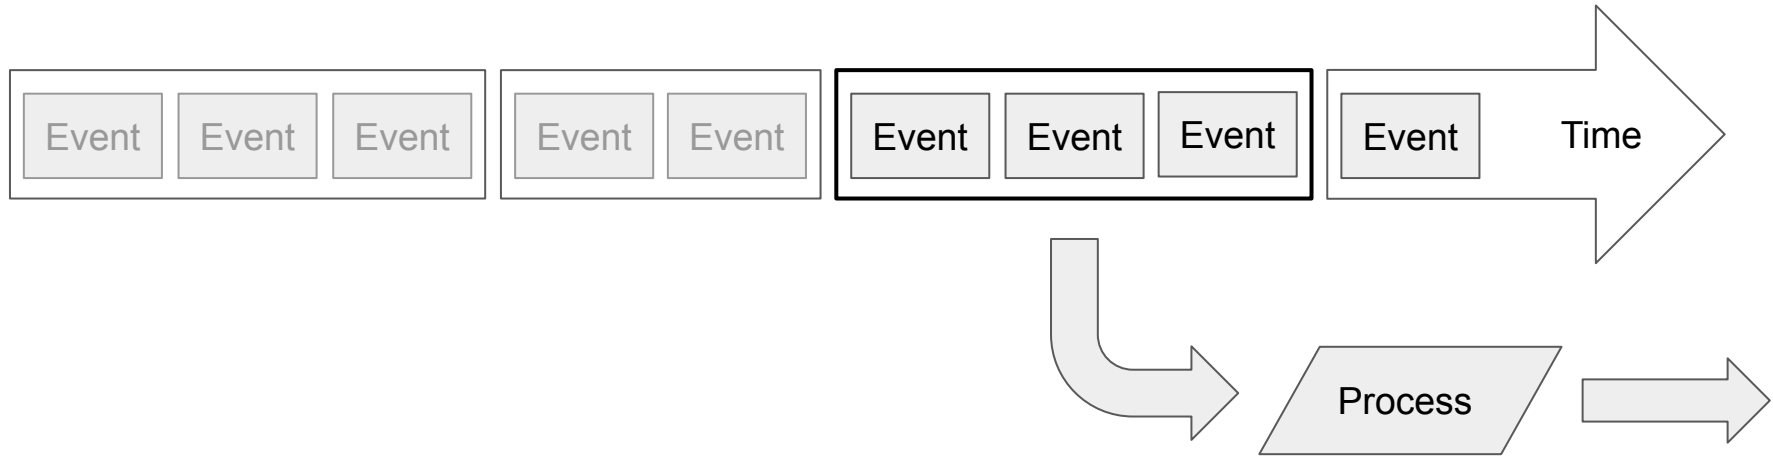
# Micro-batch approach



- ▶ Stream is cut to batch by time (for example 10 seconds batch)



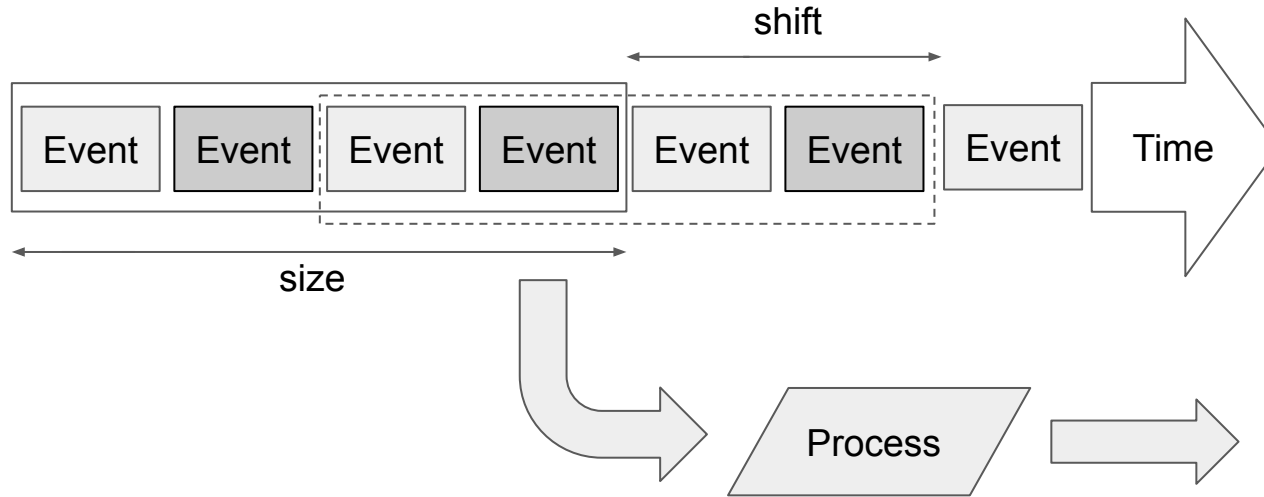
# Micro-batch approach



- ▶ Stream is cut to batch by time (for example 10 seconds batch)
- ▶ Batches are processed sequentially
- ▶ Latency  $\gg 1s$



# Windowed approach



- Batch as a sliding window
- Similar to micro-batch approach







# Event-based vs micro-batch

- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing



# Event-based vs micro-batch

- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





# Event-based vs micro-batch


- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





# Event-based vs micro-batch


- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





# Event-based vs micro-batch

- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing






# Event-based vs micro-batch



Event-based allows you to achieve less lag

Micro-batch allows you to save resources by reducing the common parts of each event handling/processing






# Event-based vs micro-batch



Event-based allows you to achieve less lag

Micro-batch allows you to save resources by reducing the common parts of each event handling/processing







# Event-based vs micro-batch



Event-based allows you to achieve less lag  
Micro-batch allows you to save resources  
by reducing the common parts  
of each event handling/processing









# Event-based vs micro-batch



Event-based allows you to achieve less lag  
Micro-batch allows you to save resources  
by reducing the common parts  
of each event handling/processing







# Event-based vs micro-batch





Event-based allows you to achieve less lag  
Micro-batch allows you to save resources  
by reducing the common parts  
of each event handling/processing





# Event-based vs micro-batch

- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing
- ▶ Micro-batch allows you to process more data on the same hardware






# Throughput vs latency

- ▶ In real world resources are restricted
- ▶ Big data needs a huge throughput => in most cases we are choosing micro-batch
- ▶ There is no right answer - you should choose the approach by task



# Throughput vs latency





## RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



## **Spark Streaming**

- ▷ Application example
- ▷ Hints




## Spark Streaming practice



## Kafka

- ▷ Internal
- ▷ CLI



- ▶ Apache Spark Streaming - classic
- ▶ Apache Spark Structured Streaming - new wave (:



Spark Streaming is an extension of Spark API core that enables real-time stream processing using micro-batch approach








# Spark Streaming





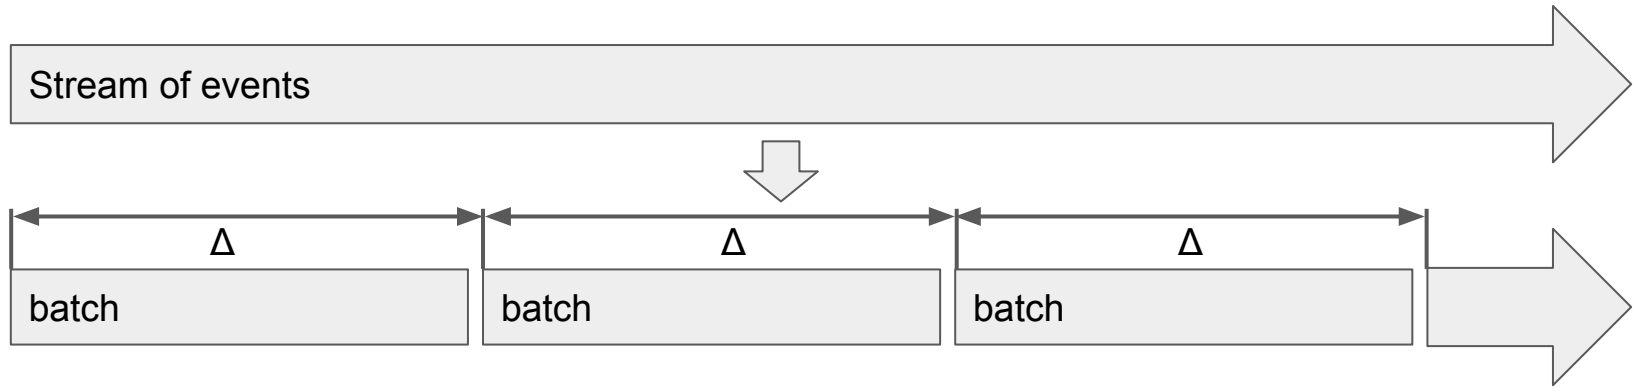
# Spark Streaming - DStream



Stream of events

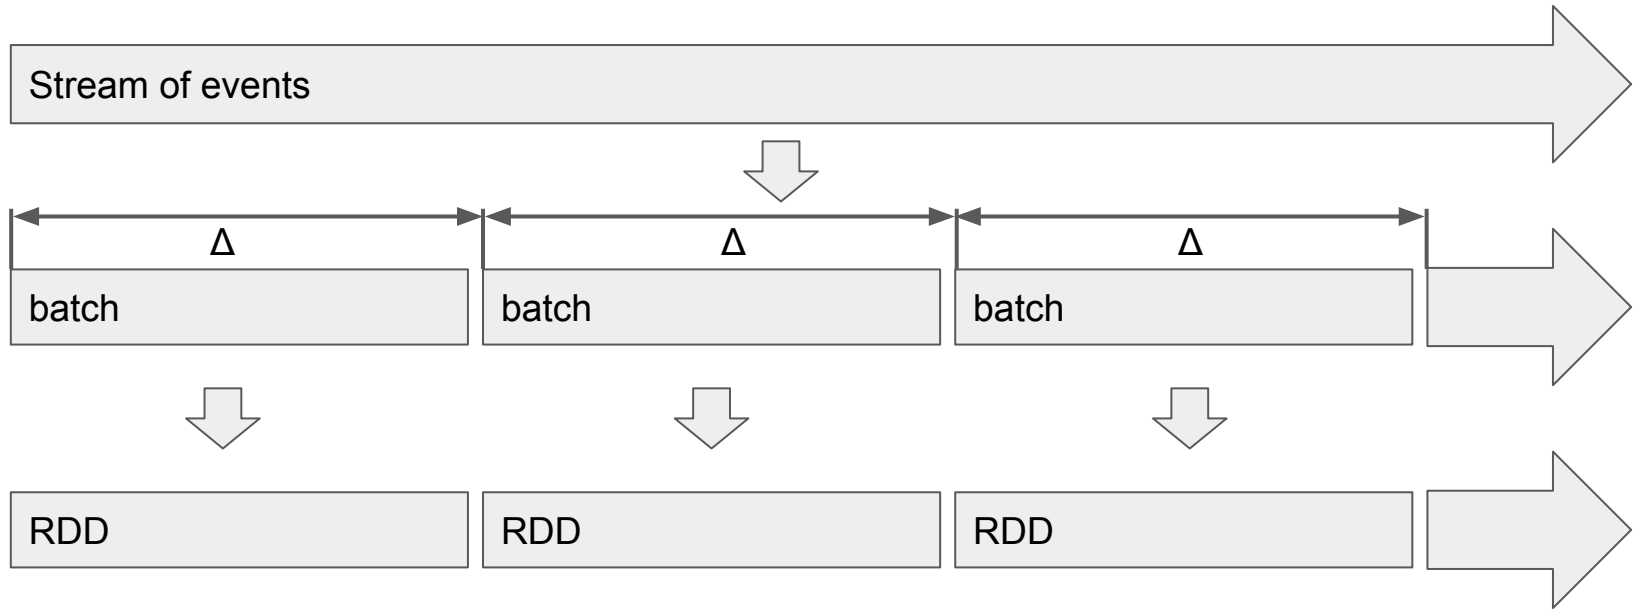


# Spark Streaming - DStream



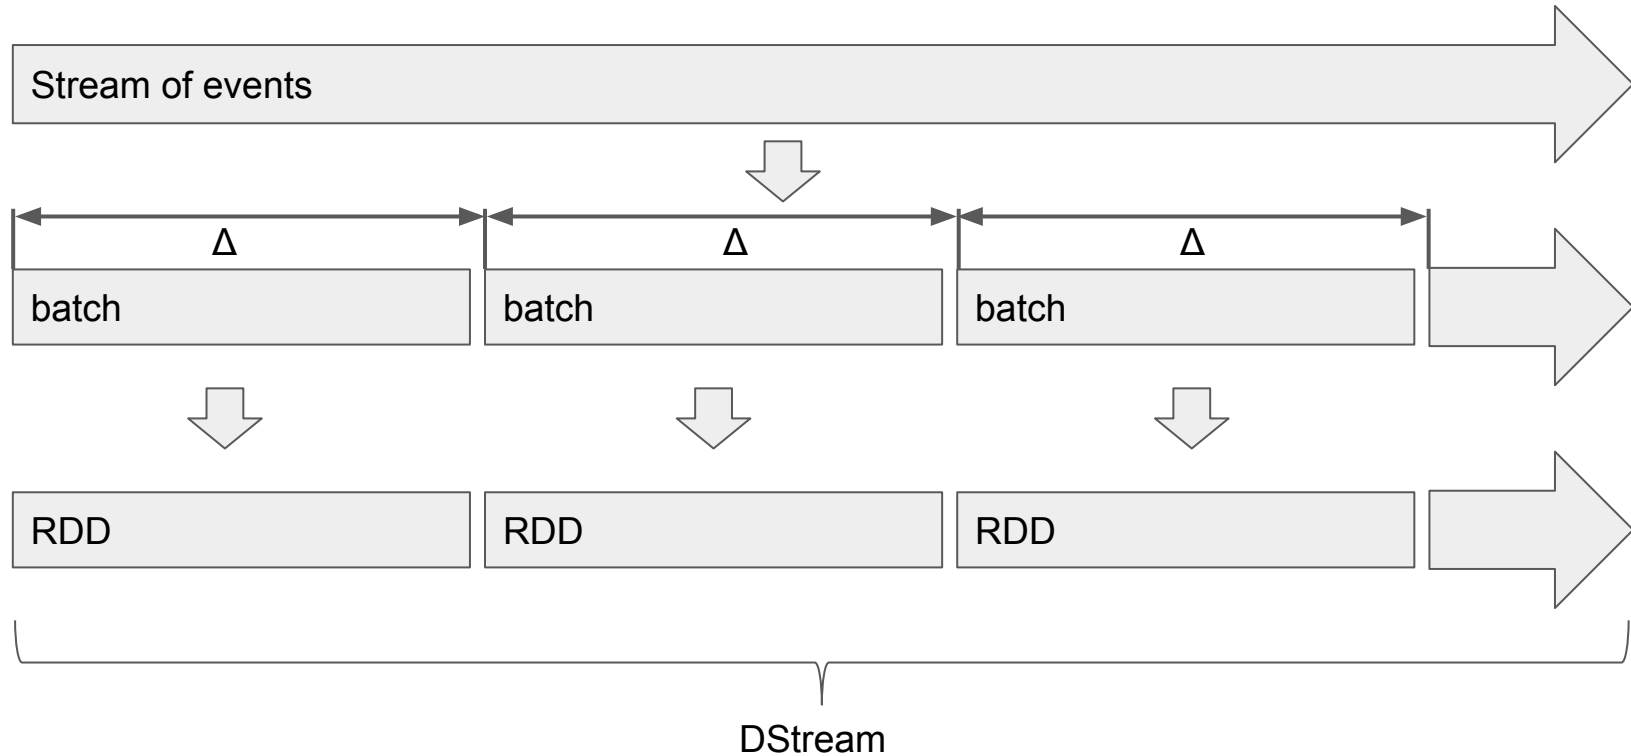


# Spark Streaming - DStream





# Spark Streaming - DStream






# Spark Streaming - processing



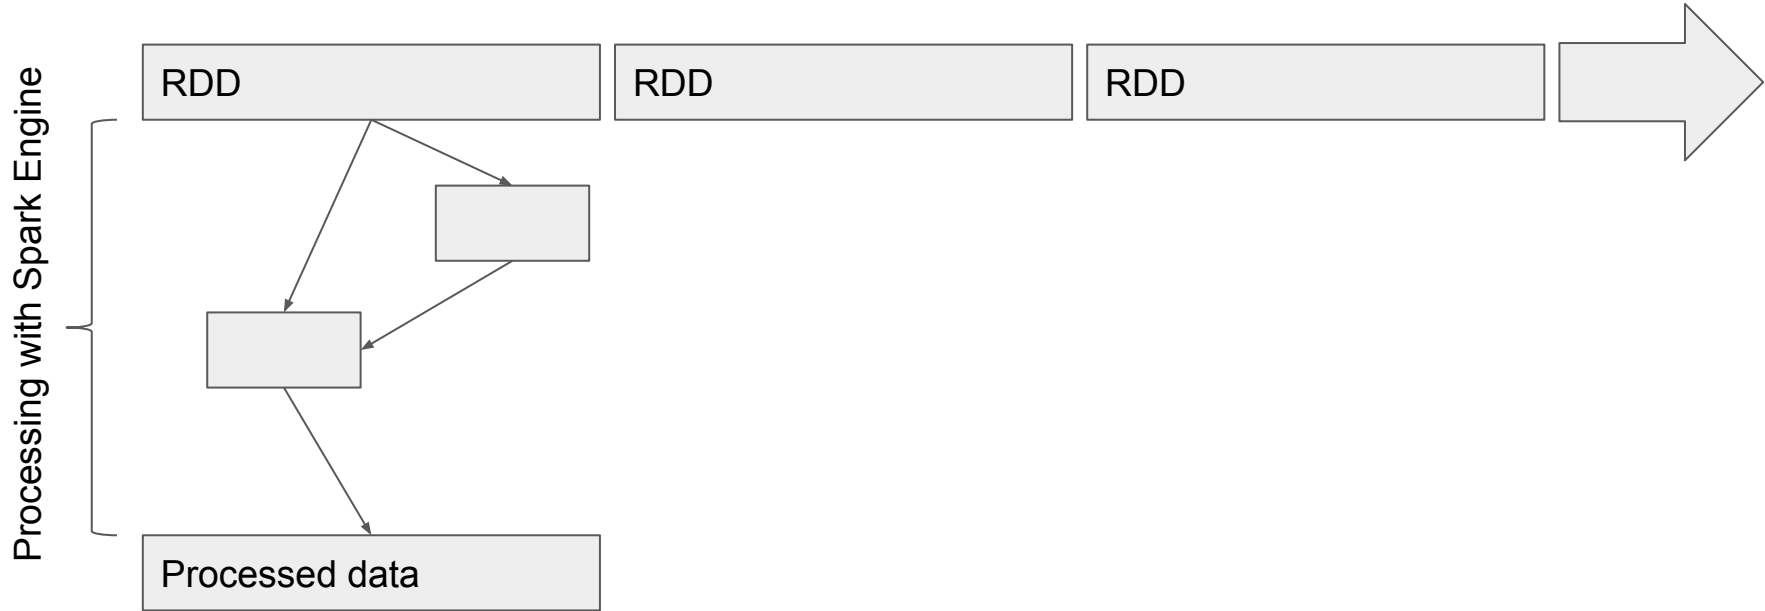


# Spark Streaming - processing





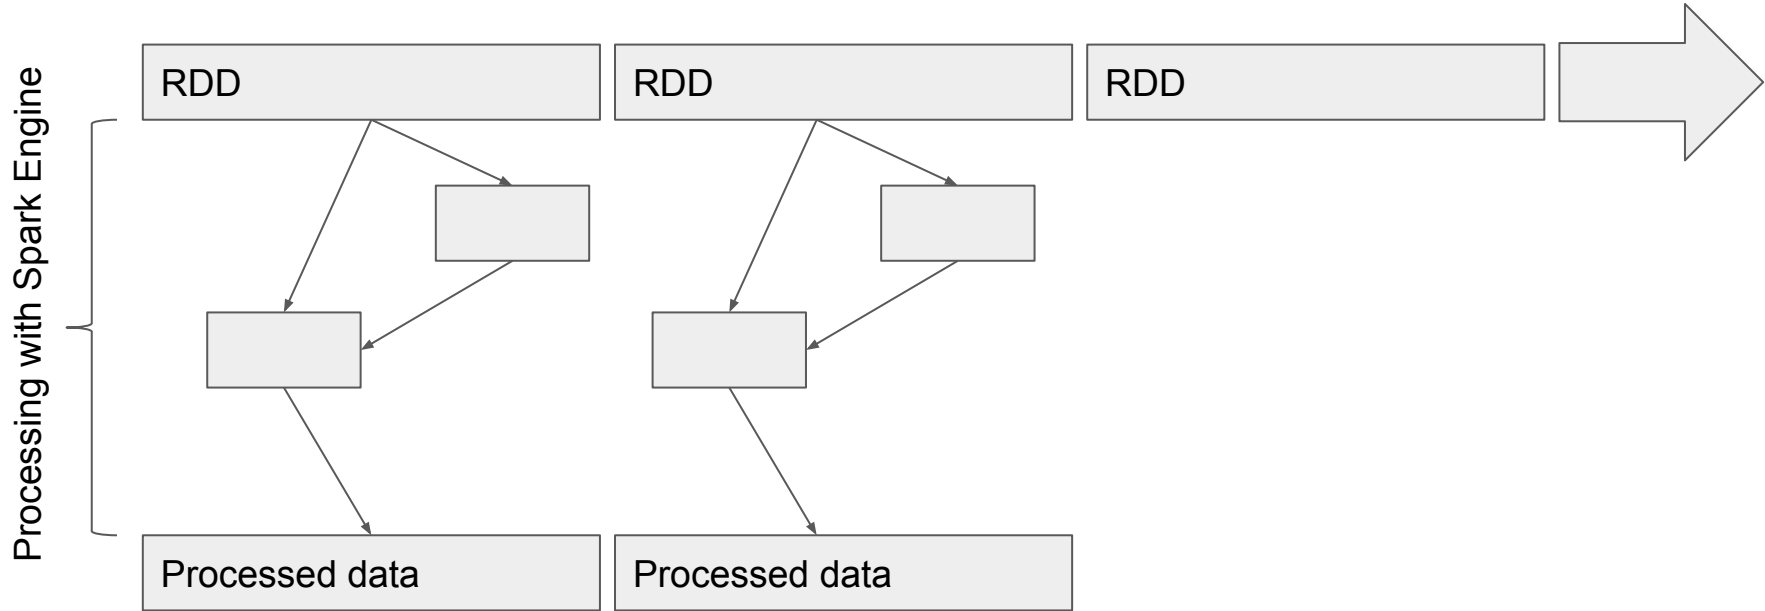
# Spark Streaming - processing





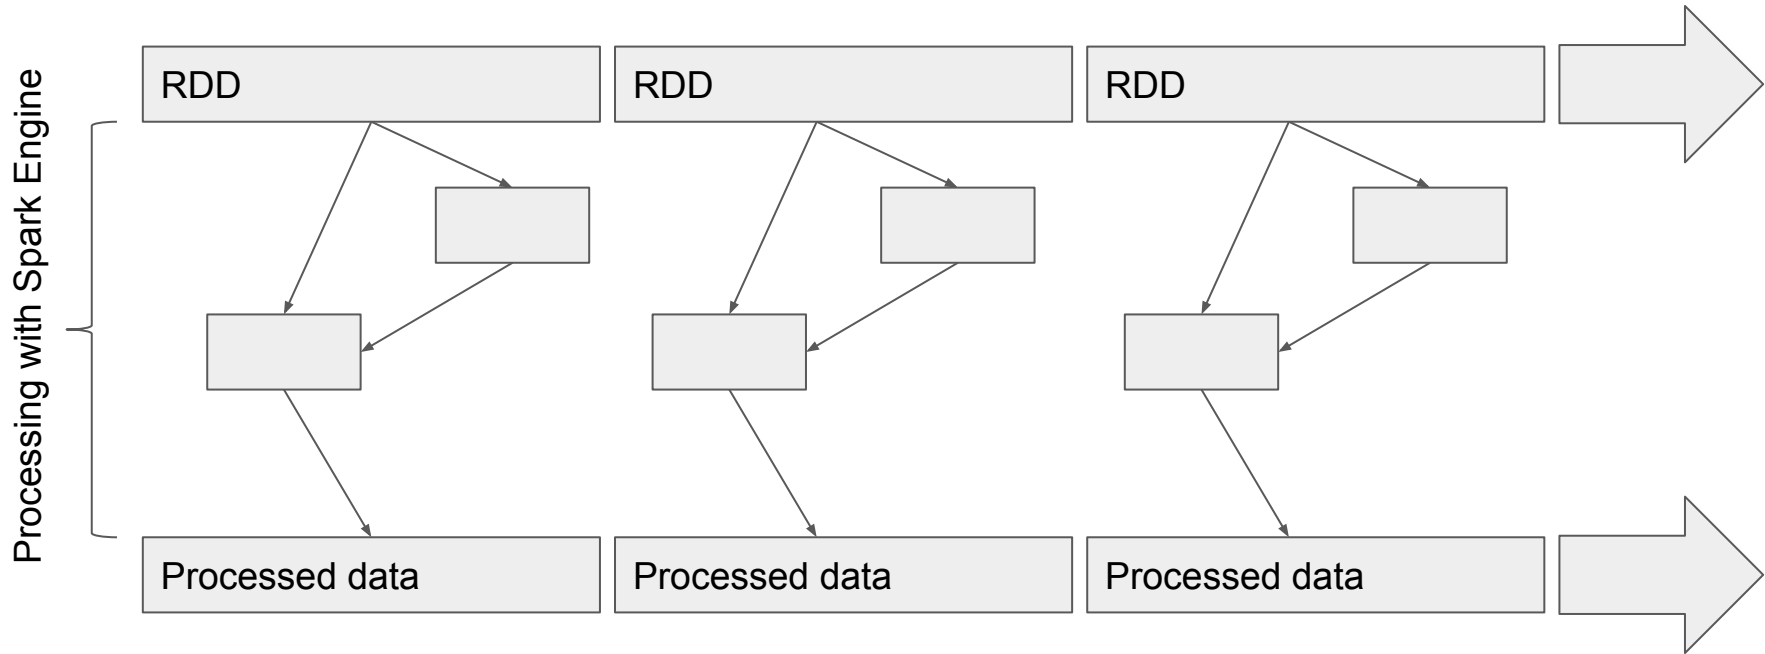


# Spark Streaming - processing





# Spark Streaming - processing






# Spark Streaming vs Spark

- ▶ Dstream has the same transformations and actions as RDD
  - ▶ *map, filter, repartition, join ...*
- ▶ Some new operations on DStream
  - ▶ *updateStateByKey(...)* - a way to make Spark Streaming stateful
  - ▶ *window(...)* - an implementation of windowed approach
  - ▶ *checkpoint(...)* - provides recovering after failures
- ▶ Spark Streaming RDD is just Spark RDD




# Stateless processing






# Stateful processing






# Stateful vs Stateless

Stateless stream  
processing



Stateful stream  
processing





# Features of Spark Streaming

- + High level abstraction
- + Processing data with Spark
- + Rich spark ecosystem (SparkSQL, SparkML, ...)



## Micro-batch

- + High throughput
- High latency



## RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



## Spark Streaming

- ▷ **Application example**
- ▷ Hints



## Spark Streaming practice



## Kafka

- ▷ Internal
- ▷ CLI





# Spark Streaming context



Spark Streaming context <- func(Spark context)

```
1 from pyspark import SparkContext
2 from pyspark.streaming import StreamingContext
3
4 sc = SparkContext(master='local[4]')
5 ssc = StreamingContext(sc, batchDuration=10)
```



# Input sources for Spark Streaming

- ▶ DStream <- func(Spark streaming context)
- ▶ Basic sources
  - ▶ `ssc.fileStream(...)` - folder with files
  - ▶ `ssc.socketTextStream(...)` - network socket
  - ▶ `ssc.queueStream(...)` - queue of RDDs for testing purposes
- ▶ Advanced sources: Kafka, Flume, Kinesis
- ▶ Custom sources (write yourself)



# Input sources examples



Create DStream over network socket

```
1 dstream = ssc.socketTextStream(hostname='localhost', port=9999)
```



Create DStream over Kafka topic

```
1 from pyspark.streaming.kafka import KafkaUtils
2
3 dstream = KafkaUtils.createDirectStream(
4     ssc,
5     topics=['test_topic'],
6     kafkaParams={'metadata.broker.list':
7                 'kafka00.example.org:9092,kafka01.example.org:9092'})
```



# Output sources for Spark Streaming



*dstream.pprint()* - prints the result to stdout



*dstream.saveAsTextFiles(...)* - saves data to external storage (e.g. HDFS)



*dstream.foreachRDD(...)* - saves data manually (e.g. external key-value storage)

- The most common option
- Providing the required semantics is entirely in the developer's jurisdiction



# Output sources examples



## Save data to HDFS

```
1 dstream.saveAsTextFiles('hdfs://cluster/path/to/result/')
```



## Save data manually

```
1 dstream.foreachRDD(lambda rdd: write_result_to_db(rdd))
```



# Sample Spark Streaming application

- ▶ Create context
- ▶ Create DStream
- ▶ Process DStream on Spark
- ▶ Write the result
- ▶ Start streaming
- ▶ Wait for exit (only for CLI)

```
1 from pyspark import SparkContext
2 from pyspark.streaming import StreamingContext
3
4 sc = SparkContext(master='local[4]')
5 ssc = StreamingContext(sc, batchDuration=10)
6
7 dstream = ssc.socketTextStream(hostname='localhost', port=9999)
8
9 result = dstream \
10     .filter(bool) \
11     .count()
12
13 result.pprint()
14
15 ssc.start()
16 ssc.awaitTermination()
```



# How to run Spark Streaming application from CLI

- ▶ Run the netcat, that can send data into socket  
`$ nc -lk 9999`
- ▶ Create *spark\_streaming\_example.py* file and put the source code in it
- ▶ Run python file with Spark  

```
1 $ spark-submit spark_streaming_example.py
```
- ▶ Send some data through netcat
- ▶ You can exit by pressing *Ctrl-C*



# How to run Spark Streaming application from iPython

- ▶ Run the netcat, that can send data into socket  
`$ nc -l -k 9999`
- ▶ Run application code (first 15 rows) in iPython
- ▶ Send some data through netcat
- ▶ You can exit by calling method  
`1 ssc.stop()`





## RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



## Spark Streaming

- ▷ Application example
- ▷ **Hints**



## Spark Streaming practice



## Kafka

- ▷ Internal
- ▷ CLI



- ▶ *dstream.repartition(num\_of\_partitions)* - distributes the data across the specified number of executors in the cluster before further processing
- ▶ Allows you to get linear scalability and data shuffle
- ▶ An example of code

```
1 dstream.repartition(30)
```



- ▶ Broadcast - creates a local copy of a read-only variable on each server
- ▶ Useful for sharing config or complex data structures between Spark Streaming stages
- ▶ An example of code

```
1 multiplier = config.get_multiplier() # will be run once
2 broadcast_multiplier = sc.broadcast(multiplier)
3
4 result = dstream\
5     .map(lambda x: float(x) * broadcast_multiplier.value)
```




# UpdateStateByKey

- ▶ `dstream.updateStateByKey(update_func, ...)` - allows you to store key-value pairs and update the values with every batch of data
- ▶ It gives the ability to obtain stateful calculations in Spark Streaming
- ▶ To use `updateStateByKey`:
  - ▶ Set the checkpoint directory
  - ▶ Specify the function how to update the state




# Example of updateStateByKey usage

Input stream



State by key








# Example of updateStateByKey usage

Input stream

Extract key, value pairs  
from input stream

Extracted pairs


State by key






# Example of updateStateByKey usage

Input stream





Extract key, value pairs  
from input stream

Extracted pairs



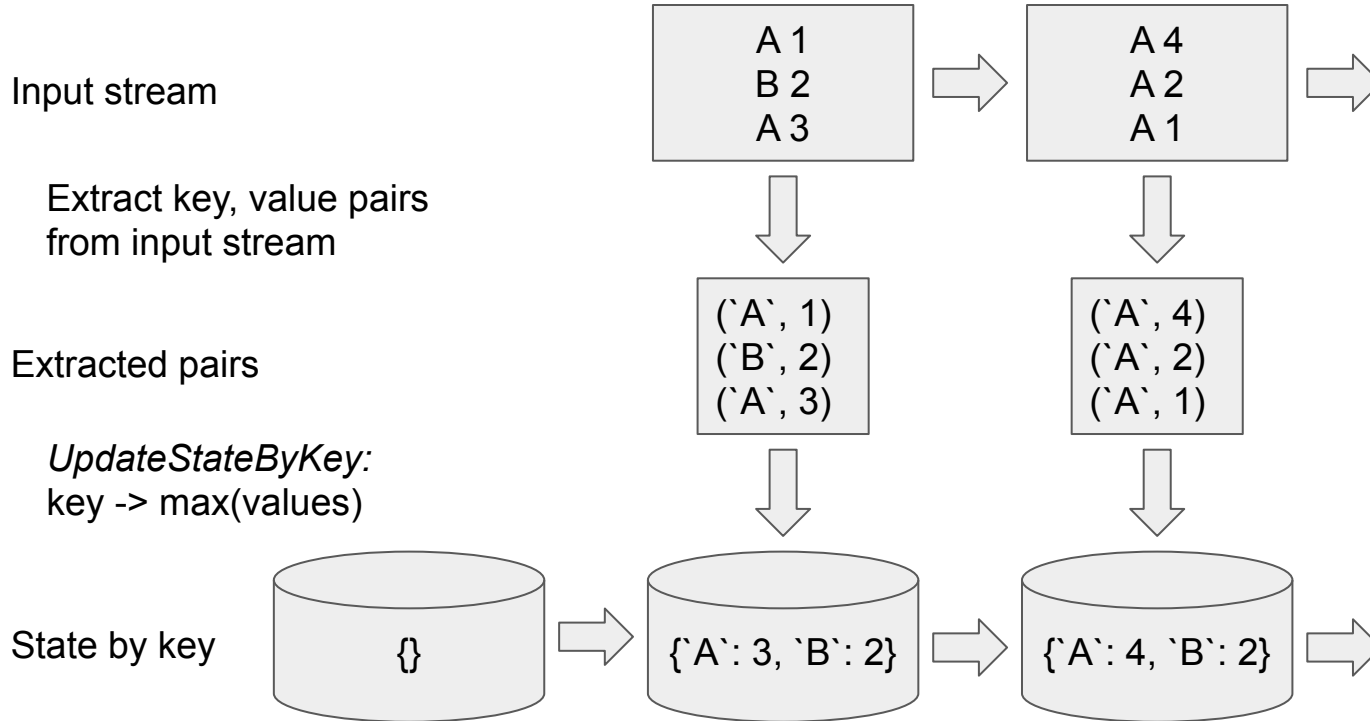
*UpdateStateByKey:*  
key -> max(values)

State by key





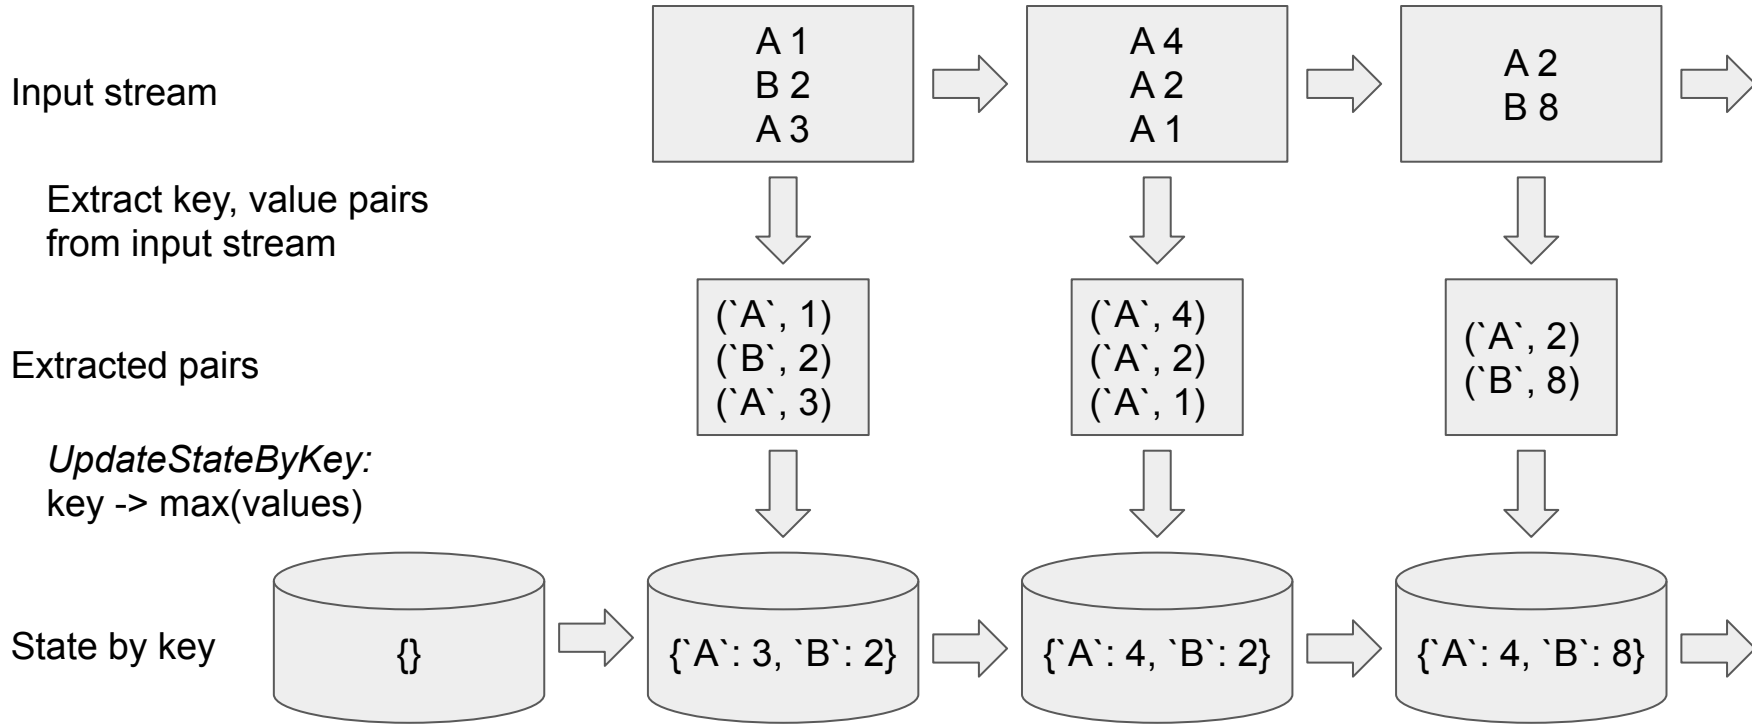
# Example of updateStateByKey usage







# Example of updateStateByKey usage





# Example of updateStateByKey usage



Find the total maximum for each key

```
1 def update_func(new_values, state):
2     max_new = max(new_values) if new_values else None
3     return max(max_new, state)
4
5 result = dstream \
6     .map(lambda line: line.split()) \
7     .map(lambda (x,y): (x, int(y))) \
8     .updateStateByKey(update_func)
```



- ▶ *dstream.window(...)* - provides a windowed approach, which allows you to apply transformations over a sliding window of data
  - ▶ Size - window duration (should be multiple of the slide duration)
  - ▶ Shift - sliding duration (should be multiple of the batch duration)
- ▶ The windowed approach is suitable to determine batches with intersection or holes between them



# Window example

- ▶ Monitoring response time of the server
- ▶ The response time comes in Spark Streaming as a stream of numbers from a network socket. It is required to determine median response time for the last minute with seconds precision

```
1 ssc = StreamingContext(sc, batchDuration=1)
2
3 dstream = ssc.socketTextStream(hostname='localhost', port=9999)
4
5 def print_median(rdd):
6     print rdd.mean() if not rdd.isEmpty() else 0
7
8 dstream\
9     .window(windowDuration=60, slideDuration=1)\
10    .map(float)\
11    .foreachRDD(print_median)
```



## RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



## Spark Streaming

- ▷ Application example
- ▷ Hints



## **Spark Streaming practice**



## Kafka

- ▷ Internal
- ▷ CLI



## RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



## Spark Streaming

- ▷ Application example
- ▷ Hints



## Spark Streaming practice




## **Kafka**

- ▷ Internal
- ▷ CLI



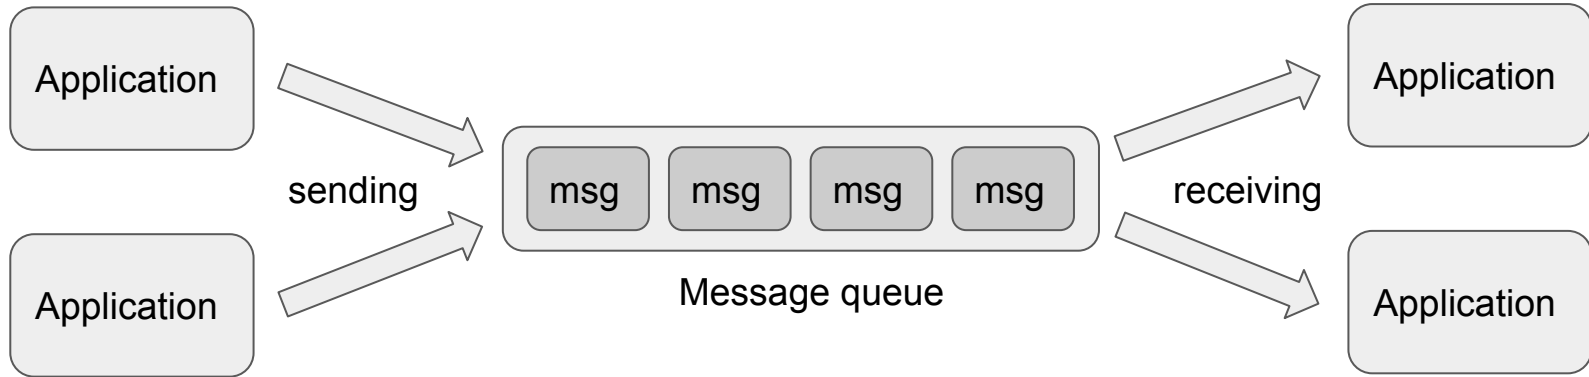
# Requirements of storage for input data



- ▶ Event stream
- ▶ Big throughput (hundreds of thousands message per second)
- ▶ Small latency (less than 1s)



# Stream handling - message queue



Message queue provides an asynchronous communications protocol between applications or between processes/threads inside a single application





# Key features of classical message queue



## Complex schemes of message delivery

- ▶ Reduces the throughput



## Per-message state

- ▶ Reduces the throughput



## Stores the data in RAM

- ▶ Not persistent storage
- ▶ Strongly limits the amount of stored data



# Storage of events in big data world

- ▶ Simplify the message delivery scheme
- ▶ Per message state not tracked by queue
- ▶ High throughput
- ▶ Persistent storage for a big amount of data




Google Cloud Pub/Sub





- ▶ Kafka is a unified, high-throughput, low-latency platform for handling real-time data feeds
- ▶ Kafka is a data bus for big data
- ▶ Kafka is an input events storage for real-time processing
- ▶ Kafka is an event-based real-time processing engine (Kafka Streams)







# Real-Time Big Data



## RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



## Spark Streaming

- ▷ Application example
- ▷ Hints



## Spark Streaming practice




## Kafka

- ▷ **Internal**
- ▷ CLI

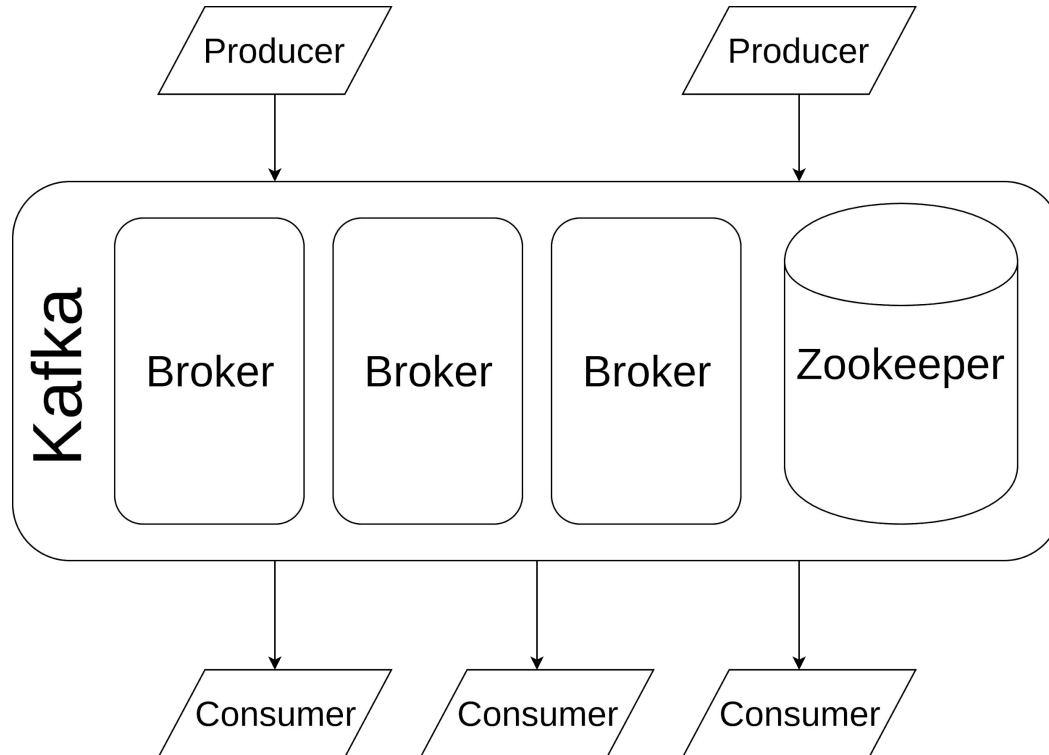


# Kafka architecture





# Kafka architecture





Topic





Partition 0

0	1	2	3
---	---	---	---

Partition 1

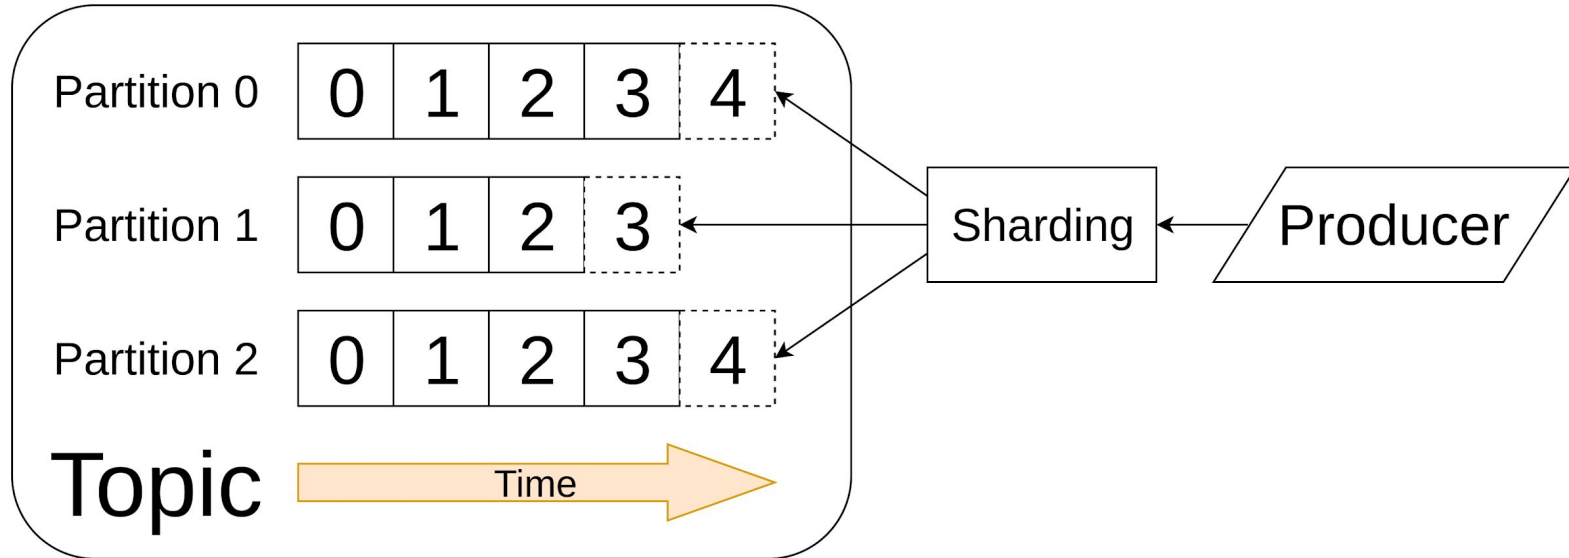
0	1	2
---	---	---

Partition 2

0	1	2	3
---	---	---	---

**Topic**




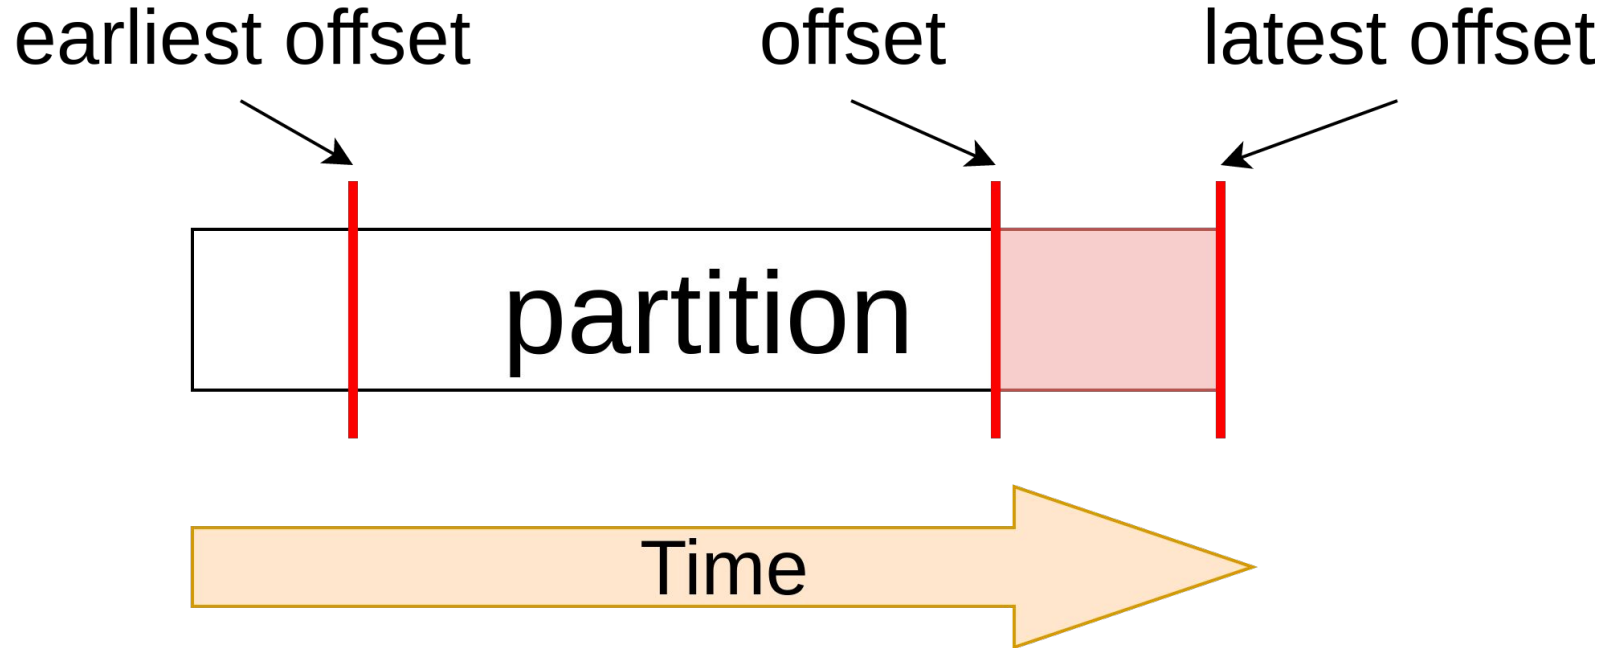






earliest offset


latest offset














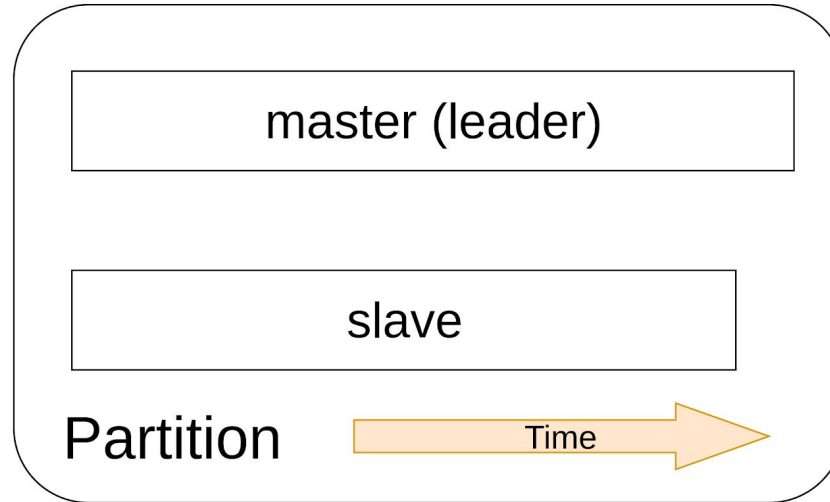
# Kafka replication





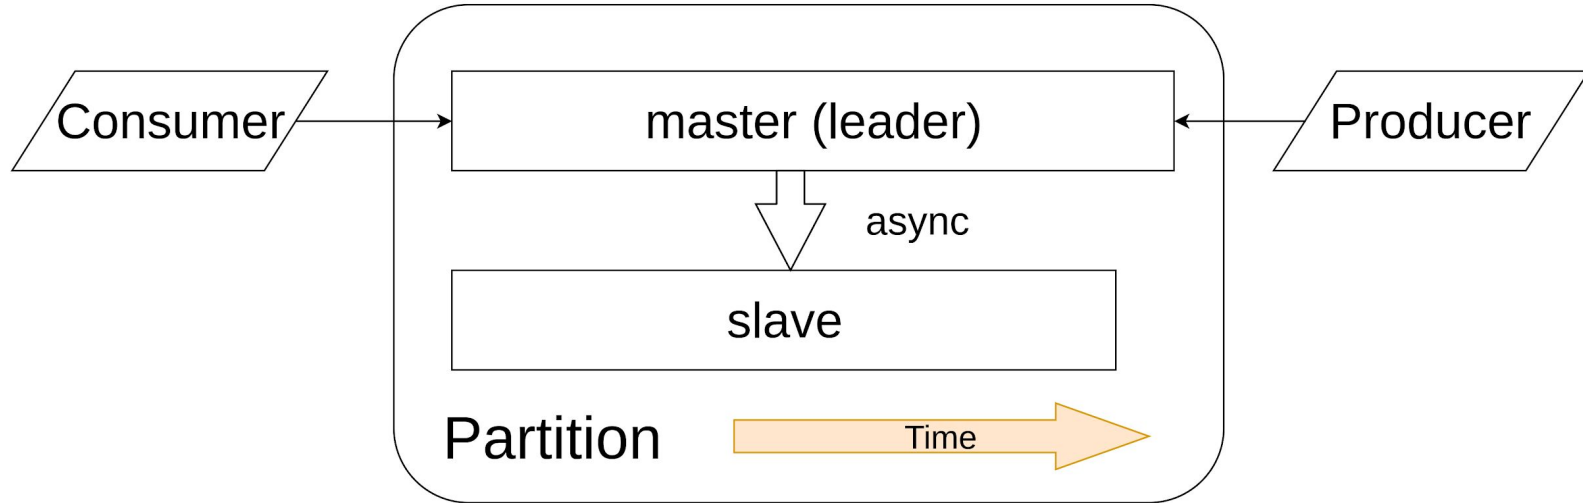


# Kafka replication






# Kafka replication





# Kafka replication delay





## RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



## Spark Streaming

- ▷ Application example
- ▷ Hints



## Spark Streaming practice



## Kafka

- ▷ Internal
- ▷ **CLI**



# Kafka CLI: kafka-topics

- ▶ It is the utility to create, delete, describe or change the topic
- ▶ `kafka-topics --zookeeper $ZOOKEEPERS --create --topic test_topic --partitions 3 --replication-factor 2`
- ▶ `kafka-topics --zookeeper $ZOOKEEPERS --describe --topic test_topic`
- ▶ `kafka-topics --zookeeper $ZOOKEEPERS --list`



# Kafka CLI: kafka-console-producer



It is utility to send data from standard input and to Kafka topic



```
kafka-console-producer --broker-list $BROKERS --topic  
test_topic
```



# Kafka CLI: kafka-console-consumer



It is utility to read data from Kafka topic



```
kafka-console-consumer --zookeeper $ZOOKEEPERS --topic  
test_topic --from-beginning
```



# Kafka CLI: kafka-run-class

- ▶ It is entry point to run any class in the Kafka environment

- ▶ `kafka-run-class kafka.tools.GetOffsetShell --broker-list $BROKERS --topic test_topic --time -1`





## RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



## Spark Streaming

- ▷ Application example
- ▷ Hints



## Spark Streaming practice



## Kafka

- ▷ Internal
- ▷ CLI



**BIGDATA  
TEAM**

# Thank you! Questions?

**Vybornov Artyom**, [avybornov@bigdatateam.org](mailto:avybornov@bigdatateam.org)

Big Data Instructor, <http://bigdatateam.org/>

Head of Big Data Dev Team, Rambler Group

<https://www.linkedin.com/in/artvybor/>